

Performance Analysis

P. K. Roy

Asst. Professor

Siliguri Institute of Technology

What is it ?

The action or process of performing a task or function.

<u>Airplane</u>	<u>Passengers</u>	<u>Range (mi)</u>	<u>Speed (mph)</u>
Boeing 737-100	101	630	598
Boeing 747	470	4150	610
BAC/Sud Concorde	132	4000	1350
Douglas DC-8-50	146	8720	544

• *Which plane is the best?*

- which metric do we use?

• *How much faster is the Concorde compared to the 747?*

Which plane has the longest route?

• *Which plane has the best transporting capacity?*

Performance Analysis

- Measure, Report, and Summarize
- Make intelligent choices
- See through the marketing hype
- Key to understanding underlying organizational motivation

Why is some hardware better than others for different programs?

*What factors of system performance are hardware related?
(e.g., Do we need a new machine, or a new operating system?)*

How does the Instruction Set Architecture (ISA) affect performance?

Computer Performance

- **Response Time or Latency**

- How long does it take for my job to run?
- How long does it take to execute a job?
- How long must I wait for the database query?

- **Throughput**

- How many jobs can the machine run at once?
- What is the average execution rate?
- How much work is getting done?

Q1: If we upgrade a machine with a new processor what do we increase?

Q2: If we add a new machine to the lab what do we increase?

Execution Time

Response (or Elapsed) Time:

Total time to complete a task including time spent executing on the CPU, accessing disk & memory, waiting for I/O & other processes, and OS overhead.

CPU Execution Time (simply CPU time):

Total time CPU spends computing on a given task (excluding the time for I/O or running other programs).

CPU time = user CPU time + System CPU execution time
= total time CPU spends only in the program execution + total time OS spends executing tasks for the program

Exercise:

A program have a system CPU time of 20 seconds, a user CPU time of 90 seconds, and a response time of 150 seconds. Calculate the followings –

- i) CPU execution time
- ii) time for I/O and other processes

Computer Performance

- For some program running on machine A,
 $\text{Performance}_X = 1 / \text{Execution time}_A$
- “A is n times faster than B”
 $\text{Performance}_A / \text{Performance}_B = n$

$$\text{Speedup} = n = \frac{\text{Performance}_A}{\text{Performance}_B} = \frac{\text{Execution Time}_B}{\text{Execution Time}_A}$$



- Problem:
 - machine A runs a program in 20 seconds
 - machine B runs the same program in 25 seconds
 - **How much faster is A compared to B?**

Clock Cycles

- Instead of reporting execution time in seconds, we often use **cycles**
- The CPU clock rate depends on the specific CPU organization (design) and hardware implementation technology (VLSI) used

$$\frac{\text{seconds}}{\text{program}} = \frac{\text{cycles}}{\text{program}} \times \frac{\text{seconds}}{\text{cycle}}$$

Clock “ticks” indicate when to start activities (one abstraction):



- cycle time = time between ticks = **seconds per cycle**
- clock rate (frequency) = **cycles per second** (1 Hz. = 1 cycle/sec)
- A 200 Mhz. clock has a cycle time $\frac{1}{200 \times 10^6} \times 10^9 = 5 \text{ nanoseconds}$

Cycles/sec = Hertz = Hz
MHz = 10^6 Hz GHz = 10^9 Hz

Nanosecond = nsec = ns = 10^{-9} second
MHz = 10^6 Hz

How to Improve Performance

So, to improve performance (everything else being equal) you can either

_____ the no. of required cycles for a program, or

_____ the clock cycle time or, said another way,

_____ the clock rate.

Fill the blanks with either increase or decrease

How to Improve Performance

- Put it another way:

$$\text{CPU time} = N_{\text{cycles}} * t_{\text{clock}} = N_{\text{cycles}} / f_{\text{clock}}$$

Where


N_{cycles} : No. of cycles

t_{clock} : cycle time or clock time (seconds per cycle)

f_{clock} : cycle rate (cycles per second)

How many cycles are required for a program?


- Could assume that no. of cycles = no. of instructions



*This assumption is incorrect,
different instructions take different amounts of time on different
machines.*

Why?

Different numbers of cycles for different instructions



- Multiplication takes more time than addition
- Floating point operations take longer than integer ones
- Accessing memory takes more time than accessing registers
- Thus, *a single machine instruction may take one or more CPU cycles to complete termed as the Cycles Per Instruction (CPI).*
 - **Instructions Per Cycle = IPC = 1/CPI**
- *Important point: changing the cycle time often changes the number of cycles required for various instructions.*

Understanding Cycles

- A given program will require
 - some number of instructions (machine instructions)
 - some number of cycles
 - some number of seconds
- We have a vocabulary that relates these quantities:
 - cycle time (seconds per cycle)
 - clock rate (cycles per second)
 - **CPI (cycles per instruction)**
a floating point intensive application might have a higher CPI
 - **MIPS (millions of instructions per second)**
this would be higher for a program using simple instructions

Performance

- Performance is determined by execution time
- Do any of the other variables equal performance?
 - No. of cycles to execute program?
 - No. of instructions in program?
 - No. of cycles per second?
 - Average no. of cycles per instruction?
 - average no. of instructions per second?
- **Common pitfall:** thinking one of the variables is indicative of performance when it really isn't.

Performance

For a specific program compiled to run on a specific machine (CPU) “A”, has the following parameters:

- *The total executed instruction count of the program.*
- *The average number of cycles per instruction (average CPI or effective CPI).*
- *Clock cycle of machine “A”*

$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

(This equation is commonly known as the CPU performance equation)

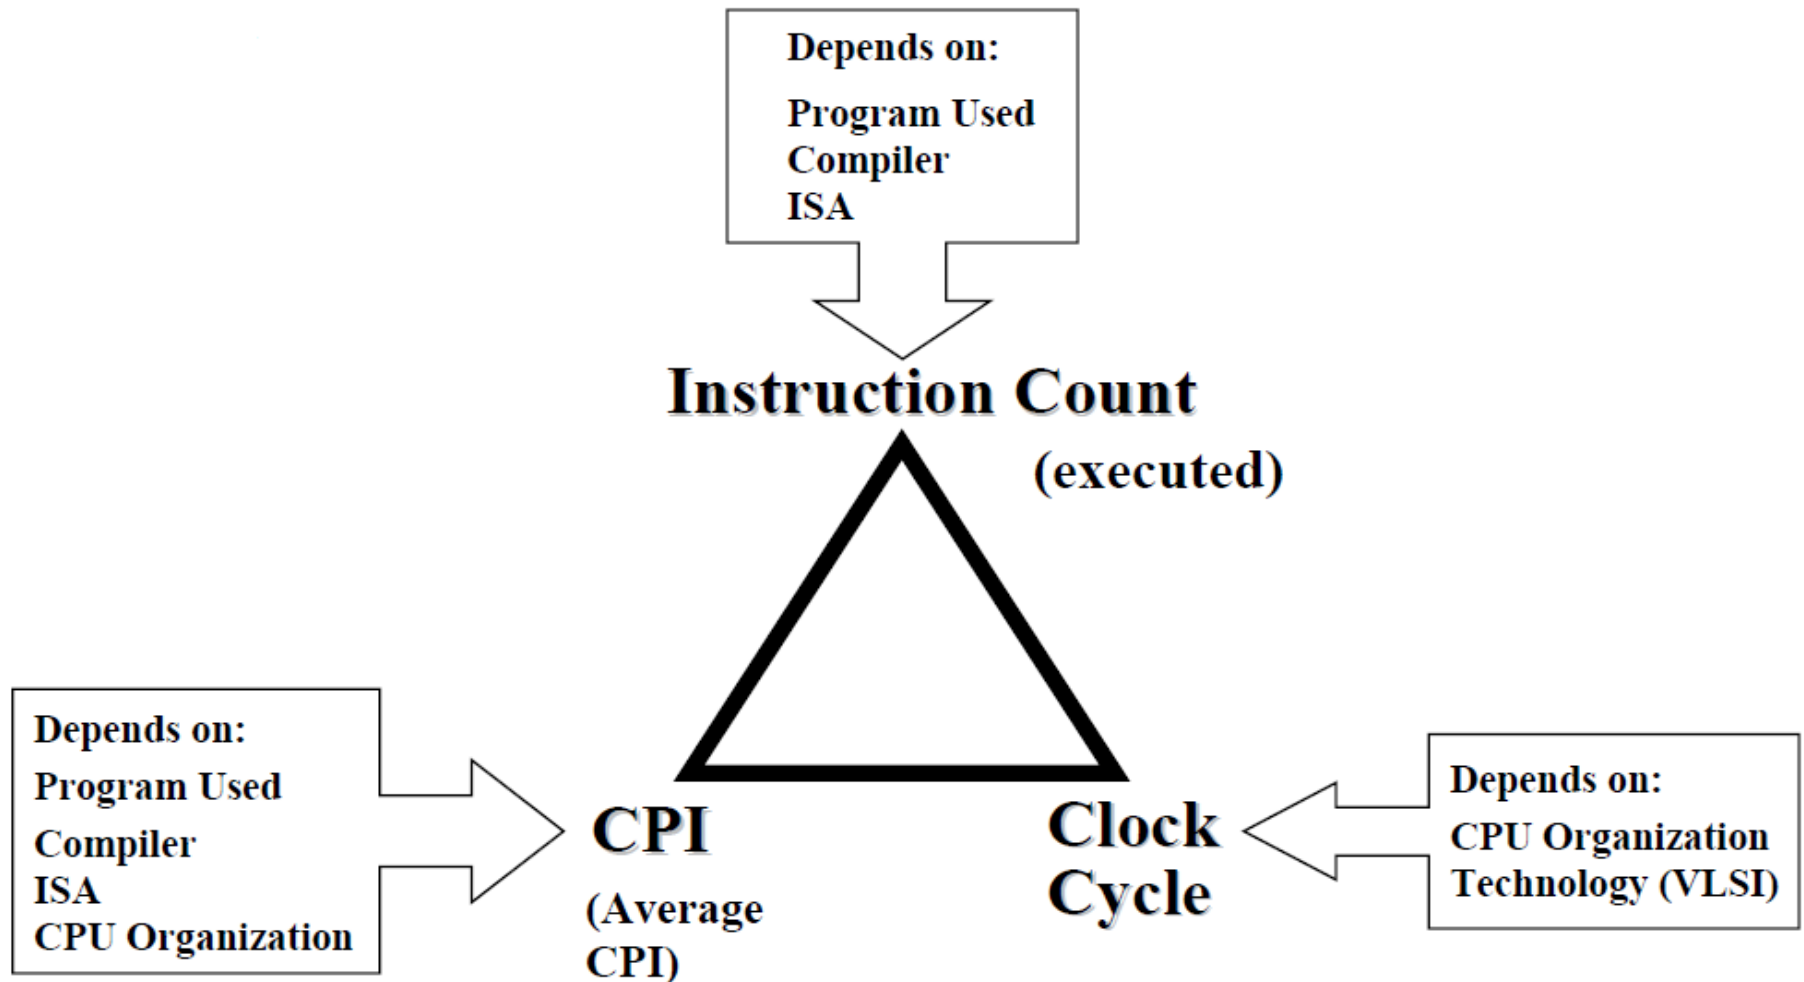
Famous equation:

$$t_{exec} = N_{instructions} \bullet CPI \bullet t_{clock}$$

note that : $N_{cycles} = N_{instructions} \bullet CPI$

CPU Execution Time

$$t_{exec} = N_{instructions} \cdot CPI \cdot t_{clock}$$



CPU Execution Time: Example

- A Program is running on a specific machine (CPU) with the following parameters:
 - Total executed instruction count: 10,000,000 instructions
 - Average CPI for the program: 2.5 cycles/instruction.
 - CPU clock rate: 200 MHz. (clock cycle = 5×10^{-9} seconds)
i.e 5 nanoseconds
- What is the execution time for this program:

CPU time	=	$\frac{\text{Seconds}}{\text{Program}}$	=	$\frac{\text{Instructions}}{\text{Program}}$	x	$\frac{\text{Cycles}}{\text{Instruction}}$	x	$\frac{\text{Seconds}}{\text{Cycle}}$
----------	---	---	---	--	---	--	---	---------------------------------------

CPU time

$$t_{exec} = N_{instructions} \cdot CPI \cdot t_{clock}$$

$$= 10,000,000 \quad \times \quad 2.5 \quad \times \quad 1 / \text{clock rate}$$

$$= 10,000,000 \quad \times \quad 2.5 \quad \times \quad 5 \times 10^{-9}$$

$$= 0.125 \text{ seconds}$$

Performance Comparison: Example

- From the previous example: A Program is running on a specific machine (CPU) with the following parameters:
 - Total executed instruction count : 10,000,000 instructions
 - Average CPI for the program: 2.5 cycles/instruction.
 - CPU clock rate: 200 MHz. Thus: $1/(200 \times 10^6) = 5 \times 10^{-9}$ seconds
- Using the same program with these changes:
 - A new compiler used: New executed instruction count : 9,500,000
New CPI: 3.0
 - Faster CPU implementation: New clock rate = 300 MHz Thus: $1/(300 \times 10^6) = 3.33 \times 10^{-9}$ seconds
- What is the speedup with the changes?

$$\text{Speedup} = \frac{\text{Old Execution Time}}{\text{New Execution Time}}$$

$$\begin{aligned}\text{Speedup} &= (10,000,000 \times 2.5 \times 5 \times 10^{-9}) / (9,500,000 \times 3 \times 3.33 \times 10^{-9}) \\ &= .125 / .095 = 1.32 \\ &\text{or } 32 \% \text{ faster after changes.}\end{aligned}$$

Computer Performance Measures :

MIPS (Million Instructions Per Second) Rating

For a specific program running on a specific CPU the MIPS rating is a measure of how many millions of instructions are executed per second:

$$\begin{aligned}\text{MIPS Rating} &= \text{Instruction count} / (\text{Execution Time} \times 10^6) \\ &= \text{Instruction count} / (\text{CPU clocks} \times \text{Cycle time} \times 10^6) \\ &= (\text{Instruction count} \times \text{Clock rate}) / (\text{Instruction count} \times \text{CPI} \times 10^6) \\ &= \text{Clock rate} / (\text{CPI} \times 10^6)\end{aligned}$$

Major problem with MIPS rating:

- MIPS rating does not account for the count of instructions executed.
- A higher MIPS rating in many cases may not mean higher performance or better execution time. i.e. due to compiler design variations.

In addition the MIPS rating:

- Does not account for the instruction set architecture (ISA) used.
- Thus it cannot be used to compare computers/CPU's with different instruction sets.

Under what conditions can the MIPS rating be used to compare performance of different CPUs?

- The MIPS rating is only valid to compare the performance of different CPUs provided that the following conditions are satisfied:
 - **The same program is used**
(actually this applies to all performance metrics)
 - **The same ISA is used**
 - **The same compiler is used**
- ⇒ (Thus the resulting programs used to run on the CPUs and obtain the MIPS rating are identical at the machine code level including the same instruction count)

MFLOPS (Million FLOating-Point Operations Per Second)

- A floating-point operation is an addition, subtraction, multiplication, or division operation applied to numbers represented by a single or a double precision floating-point representation.
- MFLOPS, for a specific program running on a specific computer, is a measure of millions of floating point-operation (megaflops) per second:

$$\text{MFLOPS} = \text{Number of floating-point operations} / (\text{Execution time} \times 10^6)$$

- MFLOPS rating is a better comparison measure between different machines (applies even if ISAs are different) than the MIPS rating.
 - Applicable even if ISAs are different
- Program-dependent: Different programs have different percentages of floating-point operations present. i.e compilers have no floating-point operations and yield a MFLOPS rating of zero.
- Dependent on the type of floating-point operations present in the program.
 - Peak MFLOPS rating for a CPU: Obtained using a program comprised entirely of the simplest floating point instructions (with the lowest CPI) for the given CPU design which does not represent real floating point programs.

Exercise

- For the multi-cycle MIPS
 - Load 5 cycles
 - Store 4 cycles
 - R-type 4 cycles
 - Branch 3 cycles
 - Jump 3 cycles
- If a program has
 - 50% R-type instructions
 - 10% load instructions
 - 20% store instructions
 - 8% branch instructions
 - 2% jump instructions

What is the CPI?

Ans: 3.6

Exercise

In a simple m/c with load-store architecture having clock rate 50 MHz, let the instruction frequency be as follows for a program –

Operations	Frequency	No. of clock cycles
ALU	40	1
Load	20	2
Store	10	2
Branch	30	2

Calculate MIPS value for the m/c.

Ans: 31.25

Aspects of CPU Performance

$$\text{CPU time} = \frac{\text{Seconds}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Cycle}}$$

	Instr. count	CPI	Clock rate
Program	X		
Compiler	X	X	
Instr. Set	X	X	
Organization	X		X
Technology		X	

Performance Enhancement Calculations: Amdahl's Law

- The performance enhancement possible due to a given design improvement is limited by the amount that the improved feature is used
- Amdahl's Law:

Performance improvement or speedup due to enhancement E:

$$\text{Speedup}(E) = \frac{\text{Execution Time without E}}{\text{Execution Time with E}} = \frac{\text{Performance with E}}{\text{Performance without E}}$$

original

- Suppose that enhancement E accelerates a fraction F of the execution time by a factor S and the remainder of the time is unaffected then:

$$\text{Execution Time with E} = ((1-F) + F/S) \times \text{Execution Time without E}$$

Hence speedup is given by:

$$\text{Speedup}(E) = \frac{\cancel{\text{Execution Time without E}}}{((1 - F) + F/S) \times \cancel{\text{Execution Time without E}}} = \frac{1}{(1 - F) + F/S}$$

F (Fraction of execution time enhanced) refers to original execution time before the enhancement is applied

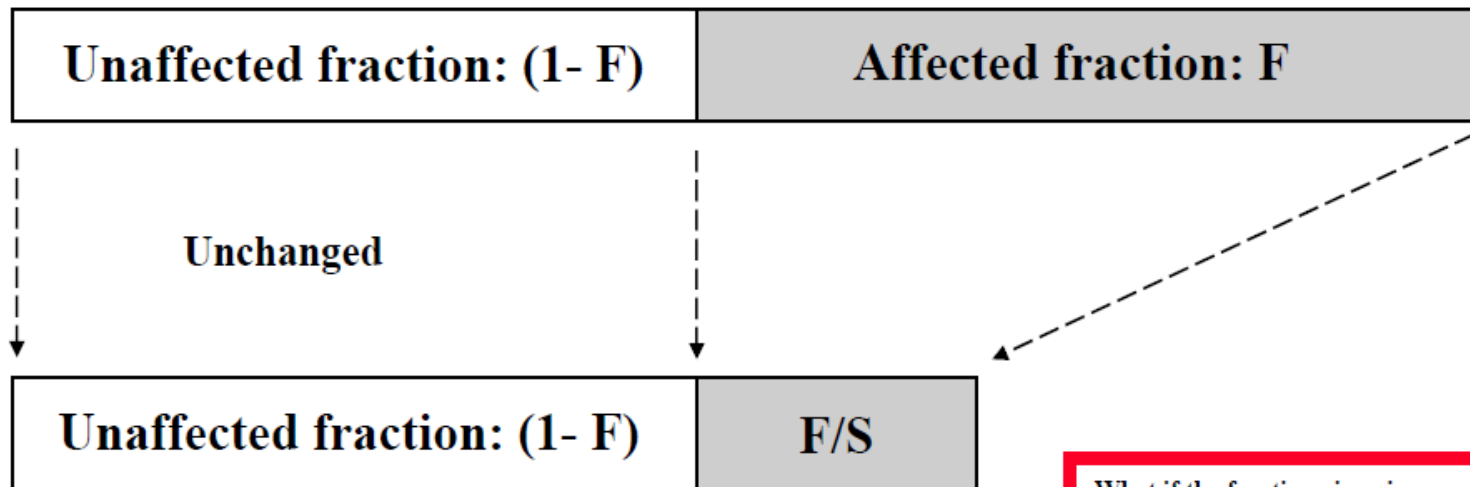
Pictorial Depiction of Amdahl's Law

Enhancement E accelerates fraction F of original execution time by a factor of S

Before:

Execution Time without enhancement E: (Before enhancement is applied)

- shown normalized to $1 = (1-F) + F = 1$



After:

Execution Time with enhancement E:

$$\text{Speedup}(E) = \frac{\text{Execution Time without enhancement E}}{\text{Execution Time with enhancement E}} = \frac{1}{(1 - F) + F/S}$$

Amdahl's Law

PARALLEL COMPUTING

Amdahl's Law: The speed-up of a program is given by

$$S_n = \frac{1}{\alpha + (1-\alpha)/n} \leq \frac{1}{\alpha} \text{ when } n = \infty$$

Where, n = number of processors and α = sequential fraction of the program

- If $\alpha = 0$, the maximum speed-up is n . However, the actual speed-up will be much less due to fixed memory size, interprocessor communication and synchronization delays.

Example

if $\alpha = 0.1$, $n = 10$

$$\text{Speedup} = \frac{1}{0.1 + \frac{0.9}{10}} \approx 5$$

As $n \longrightarrow \infty$ Speedup $\longrightarrow 10$

Observations of the Amdahl's law

- Small number of sequential operations can significantly limit speedup achievable by a parallel computer
 - This is one of the stronger arguments against parallel computers
- Amdahl's arguments serves as a way of determining whether an algorithm is a good candidate for parallelization
 - This argument doesn't take into account the problem size
 - In most applications as data size increases, the sequential part diminishes.

References

1. Advanced Computer Architecture – Kai Hwang
2. Advanced Computer Architectures – Dezso Sima, Peter Karsuk
3. Computer Architecture & Organization – John P. Hayes
4. Computer System Architecture – M. Morris Mano
5. Computer Organization & Architecture – T. K. Ghosh
6. Computer Organization & Architecture – Xpress Learning

Thank You